

MACHINE LEARNING
CREDIT APPROVAL PREDICTION

**A Project Report submitted in partial fulfillment of the requirements
for the award of the degree of**

**BACHELOR OF TECHNOLOGY
IN
COMPUTER SCIENCE AND ENGINEERING**


Submitted by

**Guntreddi Rakesh Naidu (121810305023)
V V Mani Chandra Puppala (121810305032)
H N V R Shashidhar Kappera (121810305043)
Kokkirala Vishnu (121810305008)**

Under the esteemed guidance of

Dr.J.Hyma

Associate Professor




**DEPARTMENT OF COMPUTER SCIENCE & ENGINEERING
GITAM
(Deemed to be University)
VISAKHAPATNAM
OCTOBER
2021**

**DEPARTMENT OF COMPUTER SCIENCE AND
ENGINEERING**

GITAM INSTITUTE OF TECHNOLOGY

GITAM

(Deemed to be University)



DECLARATION

I/We, hereby declare that the project report entitled “Credit Approval Prediction” is an original work done in the Department of Computer Science and Engineering, GITAM Institute of Technology, GITAM (Deemed to be University) submitted in partial fulfillment of the requirements for the award of the degree of B.Tech. in Computer Science and Engineering. The work has not been submitted to any other college or University for the award of any degree or diploma.

Date: 3-11-2021

Registration No(s).	Name(s)	Signature(s)
121810305023	Rakesh Naidu	Rakesh
121810305032	Mani Chandra	Mani
121810305043	Shashidhar	Shashidhar
121810305008	Vishnu	Vishnu

**DEPARTMENT OF COMPUTER SCIENCE AND
ENGINEERING**
GITAM INSTITUTE OF TECHNOLOGY
GITAM
(Deemed to be University)

CERTIFICATE

This is to certify that the project report entitled "**CREDIT APPROVAL PREDICTION**" is a bonafide record of work carried out by **Guntreddi Rakesh Naidu (121810305023)**, **V V Mani Chandra Puppala (121810305032)**, **H N V R Shashidhar Kapperla (121810305043)**, **Kokkirala Vishnu (121810305008)** students submitted in partial fulfillment of requirement for the award of degree of Bachelors of Technology in Computer Science and Engineering.

Project Guide

Dr.J.Hyma

Associate Professor

Head of the Department

Dr. R.Sireesha

Professor

TABLE OF CONTENTS

1	Abstract		5
2	Introduction		5
3	Literature Review		5
4	Problem Identification & Objectives		6
5	System Methodology		7
6	Overview of Technologies		7
7	Implementation		8
	7.1	Coding	8
	7.2	Testing	27
8	Results & Discussions		33
9	Conclusion and Future scope		33
10	References		33

Description

1. Abstract

Credit Card Approval Prediction is a machine learning solution to automate the process of credit card approval to an account. The model is trained with multiple attributes and the accuracy is measured with both train and test data.

2. Introduction

The accurate assessment of consumer credit risk is of uttermost importance for lending organizations. Credit scoring is a widely used technique that helps financial institutions evaluates the likelihood for a credit applicant to default on the financial obligation and decide whether to grant credit or not.

The goal here is to build an end to end automated Machine Learning solution where a user will be able to predict whether a bank customer should be approved for attaining the credit card or not. The user is only need to give the value of feature variables and the model will able to predict the binary outcome (Approve/ Not Approve). The model will be able take care of all intermediate functionalities like cross validation, hyper parameter tuning, algorithm selection etc.

This project shall be delivered in two phases:

Phase 1: All the functionalities with PyPi packages.

Phase 2: Integration of UI to all the functionalities.

Note: All the code will be written in python version 3.6

3. Literature Review

Credit card has evolved to a great level in banking industry. Each banking system consists of an enormous number of datasets to carry customer's transactions of their credit cards. So, banks would be in need of customer profiling. Customer Profiling in banks cognizes the issuer's decisions about whom to give banking facilities and what credit limit to be provided. In previous researches, profiling mainly depended on transaction data or demographic data, but in this research, both transaction and demographic data are merged in order to get more accurate results and minimize the possibility of risk occurrence

By using the best techniques, it leads to improvement in accuracy and helps banks to have high profitability through customer satisfaction by focusing on the valuable customer (companies) which are considered as the main engine in the bank's profitability. This study used k-mean, improved k-mean, fuzzy c-means and neural networks. The used dataset is labeled and for neural network classification creating a new label as a target becomes the main aspect of this study, which helps to reduce the execution time of clustering process and provide the best results with accuracy. Finally, by comparing the accuracy ratio the results show that the neural network is the best clustering technique which could give an accuracy percentage of about 98.08%.

The logistic regression (Wiginton, 1980) allows overcoming these deficiencies and became a common credit scoring tool of practitioners in financial institutions. Non-parametric techniques applied to credit scoring include the k-nearest neighbor (Henley and Hand, 1996), decision trees (Frydman et al., 1985; Davis et al., 1992), artificial neural networks (Jensen, 1992), genetic programming (Ong et al., 2005) and support vector machines (Baesens et al., 2003). More recently, research on hybrid data mining approaches has shown promising results (Lee et al., 2002; Hsieh, 2005; Lee and Chen, 2002). While the pursuit of better classifiers for credit scoring applications is a crucial research effort, improved accuracies can be easily achieved by aggregating scores predicted by an ensemble of individual classifiers. West et al. (2005) found that the accuracy of an ensemble of neural networks is superior to that of a single neural network in credit scoring and bankruptcy prediction applications. Thanks to all the researchers who published their studies.

4. Problem Identification & Objectives

The proposed project is built end to end. Starting from Data Preprocessing to Deployment. This project includes the features like:

- a) Statistical analysis
- b) Hyper parameter tuning
- c) Best algorithm selection
- d) Deployment in Heroku using flask.


The main objectives of the proposed project are to:

- a) Increase the accuracy
- b) Do Exploratory data analysis
- c) Test the model with different algorithms
- d) Try different model selection criteria
- e) Do Hyperparameter tuning
- f) Deploy the project for easy use

Different technologies or libraries used in the project are:

- a) numpy, pandas
- b) matplotlib, seaborn
- c) scipy, scikit learn
- d) xgboost
- e) html, css
- f) flask, gunicorn

5. System Methodology



6. Overview of Technologies

6.1. Numpy

NumPy is the fundamental package for scientific computing in Python. It is a Python library that provides a multidimensional array object, various derived objects (such as masked arrays and matrices), and an assortment of routines for fast operations on arrays, including mathematical, logical, shape manipulation, sorting, selecting, I/O, discrete Fourier transforms, basic linear algebra, basic statistical operations, random simulation and much more.

6.2. Pandas

Pandas is an open source Python package that is most widely used for data science/data analysis and machine learning tasks. It is built on top of another package named Numpy, which provides support for multi-dimensional arrays. As one of the most popular data wrangling packages, Pandas works well with many other data science modules inside the Python ecosystem, and is typically included in every Python distribution, from those that come with your operating system to commercial vendor distributions like ActiveState's ActivePython.

6.3. Matplotlib

Matplotlib is a cross-platform, data visualization and graphical plotting library for Python and its numerical extension NumPy. As such, it offers a viable open source alternative to MATLAB. Developers can also use matplotlib's APIs (Application Programming Interfaces) to embed plots in GUI applications.

6.4. Seaborn

Seaborn is an open-source Python library built on top of matplotlib. It is used for data visualization and exploratory data analysis. Seaborn works easily with dataframes and the Pandas library. The graphs created can also be customized easily. Below are a few benefits of Data Visualization.

6.5. Scipy

SciPy is a scientific computation library that uses NumPy underneath. SciPy stands for Scientific Python. It provides more utility functions for optimization, stats and signal processing. Like NumPy, SciPy is open source so we can use it freely. SciPy was created by NumPy's creator Travis Olliphant.

6.6. Scikit learn

Scikit-learn is a library in Python that provides many unsupervised and supervised learning algorithms. It's built upon some of the technology you might already be familiar with, like NumPy, pandas, and Matplotlib!

The functionality that scikit-learn provides include:

- ✓ **Regression**, including Linear and Logistic Regression
- ✓ **Classification**, including K-Nearest Neighbors
- ✓ **Clustering**, including K-Means and K-Means++
- ✓ **Model selection**
- ✓ **Preprocessing**, including Min-Max Normalization

6.7. Xgboost

XGBoost is a tree based ensemble machine learning algorithm which is a scalable machine learning system for tree boosting. XGBoost stands for Extreme Gradient Boosting. It uses more accurate approximations to find the best tree model. In prediction problems involving unstructured data (images, text, etc.) artificial neural networks tend to outperform all other algorithms or frameworks. However, when it comes to small-to-medium structured/tabular data, decision tree based algorithms are considered best-in-class right now.

6.8. HTML, CSS

HTML (the Hypertext Markup Language) and CSS (Cascading Style Sheets) are two of the core technologies for building Web pages. HTML provides the structure of the page, CSS the (visual and aural) layout, for a variety of devices. Along with graphics and scripting, HTML and CSS are the basis of building Web pages and Web Applications. Learn more below about:

6.9. Flask, gunicorn

Flask is a Python-based microframework that is popular with web developers, given its lightweight nature and ease of use. Gunicorn is a Python WSGI HTTP Server that

uses a pre-fork worker model. By using gunicorn, we will be able to serve your Flask application on more than one thread.

7. Implementation

7.1 Coding

Feature Engineering

The screenshot shows a Jupyter Notebook interface with the following details:

- Title:** 1. Feature_Engineering.ipynb
- Section:** Credit Approval
- Text:** Commercial banks receive a lot of applications for credit cards. Many of them get rejected for many reasons, like high loan balances, low income levels, or too many inquiries on an individual's credit report, for example. Manually analyzing these applications is mundane, error-prone, and time-consuming (and time is money). Luckily, this task can be automated with the power of machine learning and pretty much every commercial bank does so nowadays. In this notebook, we will build an automatic credit card approval predictor using machine learning techniques, just like the real banks do.
- Code:**

```
1. Gender : num 1 1 0 0 0 0 1 0 0 0 ...
2. Age : chr "58.67" "24.50" "27.83" "20.17" ...
3. Debt : num 4.46 0.5 1.54 5.62 4 ...
4. Married : chr "u" "u" "u" "u" ...
5. BankCustomer : chr "g" "g" "g" "g" ...
6. EducationLevel : chr "q" "q" "w" "w" ...
7. Ethnicity : chr "h" "h" "v" "v" ...
8. YearsEmployed : num 3.04 1.5 3.75 1.71 2.5 ...
9. PriorDefault : When you accept a credit card, you agree to certain terms. For example, you agree to make your minimum payment by the due date listed on your credit card statement. If you miss the minimum payment six months in a row, your credit card will be in default. Your credit card issuer will likely close your account and report the default to the credit bureaus.
In the months leading up to a default, your (late) payment status will be reported to the three major credit bureaus, and your credit score will be impacted by the lateness of your payments. If you apply for any new credit cards or loans after a default, your application will likely be denied because creditors think you are at risk of defaulting on any new credit obligations. In fact, some lenders will not approve you at all until you have cleared up the default balance (or it drops off your credit report).
0:default, 1:no prior default
10. Employed : num 1 0 1 0 0 0 0 0 0 ...
11. CreditScore : Lenders use credit scores to evaluate the probability that an individual will repay loans in a timely manner.
12. DriversLicense: chr "f" "f" "t" "f" ...
13. Citizen : chr "g" "g" "g" "s" ...
14. ZipCode : chr "00043" "00280" "00100" "00120" ...
15. Income : num 560 824 3 0 0 ...
16. Approved : chr "+" "+" "+" "+" ...
Credit card being held in hand
```
- Text:** We'll use the Credit Card Approval dataset from the UCI Machine Learning Repository. The structure of this notebook is as follows:
- Text:** First, we will start off by loading and viewing the dataset. We will see that the dataset has a mixture of both numerical and non-numerical features, that it contains values from different ranges, plus that it contains a number of missing entries. We will have to preprocess the dataset to ensure the machine learning model we choose can make good predictions. After our data is in good shape, we will do some exploratory data analysis to build our intuitions. Finally, we will build a machine learning model that can predict if an individual's application for a credit card will be accepted. First, loading and viewing the dataset. We find that since this data is confidential, the contributor of the dataset has anonymized the feature names.
- Code:**

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

## Display all the columns of the dataframe
pd.pandas.set_option('display.max_columns',None)

[2]: import warnings
warnings.filterwarnings('ignore')

[3]: dataset = pd.read_csv('crx.data',header=None)
dataset.head()
```
- Data Preview:**

	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
0	b	30.83	0.000	u	g	w	v	125	t	t	1	f	g	00202	0	+
1	a	59.63	4.450	w	m	m	m	204	t	t	1	f	e	00043	550	-
- Bottom Status:** Saving completed, Mode: Command, Ln 1, Col 1, 1. Feature_Engineering.ipynb

```
[4]: dataset.columns
```

```
[4]: Int64Index([0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15], dtype='int64')
```

The output may appear a bit confusing at its first sight, but let's try to figure out the most important features of a credit card application. The features of this dataset have been anonymized to protect the privacy, but Analysis of Credit Approval Data blog gives us a pretty good overview of the probable features. The probable features in a typical credit card application are Gender, Age, Debt, Married, BankCustomer, EducationLevel, Ethnicity, YearsEmployed, PriorDefault, Employed, CreditScore, DriversLicense, Citizen, ZipCode, Income and finally the ApprovalStatus. This gives us a pretty good starting point, and we can map these features with respect to the columns in the output.

As we can see from our first glance at the data, the dataset has a mixture of numerical and non-numerical features. This can be fixed with some preprocessing, but before we do that, let's learn about the dataset a bit more to see if there are other dataset issues that need to be fixed.

```
[5]: columns = ["Gender", "Age", "Debt", "Married", "BankCustomer", "EducationLevel", "Ethnicity", "YearsEmployed", "PriorDefault", "Employed", "CreditScore", "DriversLicense", "Citizen", "ZipCode", "Income", "Approved"]
```

```
[6]: dataset.columns = columns
```

```
[7]: dataset.head()
```

	Gender	Age	Debt	Married	BankCustomer	EducationLevel	Ethnicity	YearsEmployed	PriorDefault	Employed	CreditScore	DriversLicense	Citizen	ZipCode	Income	Approved
0	b	30.83	0.000	u	g	w	v	1.25	t	t	1	f	g	00202	0	+
1	a	58.67	4.460	u	g	q	h	3.04	t	t	6	f	g	00043	560	+
2	a	24.50	0.500	u	g	q	h	1.50	t	f	0	f	g	00280	824	+
3	b	27.83	1.540	u	g	w	v	3.75	t	t	5	t	g	00100	3	+
4	b	20.17	5.625	u	g	w	v	1.71	t	f	0	f	s	00120	0	+

replacing ? with nan

```
[13]: dataset.replace('?', np.nan, inplace=True)
```

```
[14]: dataset.isnull().sum()
```

	Gender	Age	Debt	Married	BankCustomer	EducationLevel	Ethnicity	YearsEmployed	PriorDefault	Employed	CreditScore	DriversLicense	Citizen	ZipCode	Income	Approved
Gender	12															
Age	12															
Debt	0															
Married	6															
BankCustomer	6															
EducationLevel	9															
Ethnicity	9															
YearsEmployed	0															
PriorDefault	0															
Employed	0															
CreditScore	0															
DriversLicense	0															
Citizen	0															
ZipCode	13															
Income	0															
Approved	0															

```
[15]: dataset['Gender'].value_counts()
```

```
[15]: b    468
```

```
[19]: ## Here we will check the percentage of nan values present in each feature
## 1 - step make the list of features which has missing values
features_with_na=[features for features in dataset.columns if dataset[features].isnull().sum()>1]
## 2- step print the feature name and the percentage of missing values

for feature in features_with_na:
    print(feature, np.round(dataset[feature].isnull().mean(), 4), ' % missing values')

Gender 0.0174 % missing values
Age 0.0174 % missing values
Married 0.0087 % missing values
BankCustomer 0.0087 % missing values
EducationLevel 0.013 % missing values
Ethnicity 0.013 % missing values
ZipCode 0.0188 % missing values
```

We replaced all the question marks with NaNs. This is going to help us in the next missing value treatment that we are going to perform.


An important question that gets raised here is why are we giving so much importance to missing values? Can't they be just ignored? Ignoring missing values can affect the performance of a machine learning model heavily. While ignoring the missing values our machine learning model may miss out on information about the dataset that may be useful for its training. Then, there are many models which cannot handle missing values implicitly.

So, to avoid this problem, we are going to impute the missing values with a strategy called mean imputation.

```
[20]: dataset['Age']=dataset['Age'].astype(float)
```

```
[21]: dataset.info()
```

```
[36]: for feature in discrete_feature:
    data=dataset.copy()
    plt.figure(figsize=(15,8))
    sns.set_style('whitegrid')
    sns.countplot(x='Approved',hue=feature,data=data)
```



```
[63]: def detect_outliers(data):
    outliers=[]
    threshold=3
    mean = np.mean(data)
    std = np.std(data)
```

```
for i in range(len(data)):
    z_score= (data[i] - mean)/std
    if np.abs(z_score) > threshold:
        outliers.append(i)
return outliers
```

```
[64]: outlier_pt=detect_outliers(data['Age'])
print(outlier_pt)
[130, 157, 206, 296, 405, 485, 539, 550, 585]
[65]: outlier_pt1=detect_outliers(data['Debt'])
print(outlier_pt1)
[44, 69, 234, 241, 250, 305, 317, 529, 550, 586]
[66]: outlier_pt2=detect_outliers(data['Income'])
print(outlier_pt2)
[]
```

handling 0 values

```
[79]: dataset[dataset['Income']==0].shape[0]
```

```
[79]: 295
```

```
[80]: dataset[dataset['CreditScore']==0].shape[0]
```

```
[80]: 395
```

```
[81]: dataset["CreditScore"].nunique()
```

```
[81]: 23
```

```
[82]: data['Income']=np.where(data['Income']==0,data['Income'].mean(),data['Income'])
```

```
[82]: Gender_b Married_I Married_u Married_y BankCustomer_g BankCustomer_gg BankCustomer_p PriorDefault_t Employed_t DriversLicense_t Citizen_p Citizen_s Age Debt EducationLevel Ethnicity YearsEmployed Cre
0 1 0 1 0 1 0 0 1 1 0 0 0 30.83 0.000 0.515625 0.423559 1.25
1 0 0 0 1 0 1 0 0 1 1 0 0 58.67 4.460 0.653846 0.630435 3.04
2 0 0 0 1 0 1 0 0 1 0 0 0 24.50 0.500 0.653846 0.630435 1.50
3 1 0 1 0 1 0 0 0 1 1 1 0 27.83 1.540 0.515625 0.423559 3.75
4 1 0 1 0 1 0 0 1 0 0 0 0 20.17 5.625 0.515625 0.423559 1.71
```

```
[83]: data.to_csv('feature-engineering.csv', index = False)
```

Feature Selection

Univariate Selection

```
[8]: from sklearn.feature_selection import SelectKBest
      from sklearn.feature_selection import chi2

[9]: ### Apply SelectKBest Algorithm
      ordered_rank_features=SelectKBest(score_func=chi2,k=15)
      ordered_feature=ordered_rank_features.fit(X,y)

[10]: dfscores=pd.DataFrame(ordered_feature.scores_,columns=["Score"])
      dfcolumns=pd.DataFrame(X.columns)

[11]: features_rank=pd.concat([dfcolumns,dfscores],axis=1)

[12]: features_rank.columns=['Features','Score']
      features_rank

[12]:
```

	Features	Score
0	Gender_b	0.175810
1	Married_J	2.495114
2	Married_u	4.908390
3	Married_y	18.817889
4	BankCustomer_g	4.908390
5	BankCustomer_gg	2.495114
6	BankCustomer_p	18.817889
7	PriorDefault_t	170.746388

Feature Importance

```
[14]: from sklearn.ensemble import ExtraTreesClassifier
      import matplotlib.pyplot as plt
      model=ExtraTreesClassifier()
      model.fit(X,y)

[14]: ExtraTreesClassifier()

[15]: print(model.feature_importances_)

[0.01796896 0.00223785 0.00678839 0.00625776 0.00557448 0.00267097
0.00763106 0.35663928 0.07549839 0.02085268 0.00599864 0.01473194
0.06499268 0.06829546 0.08091447 0.04091068 0.07608249 0.07574094
0.07021291]

[16]: ranked_features=pd.Series(model.feature_importances_,index=X.columns)
      lst=ranked_features.nlargest(10).plot(kind='barh')
      plt.show()
      lst=ranked_features.nlargest(6)
      print(lst)
```



Multicollinearity

- variance inflation factor
- correlation

variance inflation factor

```
[17]: from statsmodels.stats.outliers_influence import variance_inflation_factor
vif = pd.DataFrame()
vif["VIF Factor"]=[variance_inflation_factor(X.values,i) for i in range(X.shape[1])]
vif["features"]=X.columns
#Let's check the values
vif
```

	VIF Factor	features
0	3.412023	Gender_b
1	inf	Married_J
2	inf	Married_u
3	inf	Married_y
4	inf	BankCustomer_g
5	inf	BankCustomer_gg
6	inf	BankCustomer_p
7	3.360099	PriorDefault_t
8	3.853341	Employed_t
9	1.952504	DriversLicense_t

correlation

```
[18]: import seaborn as sns
corr=df.iloc[:, :-1].corr()
top_features=corr.index
plt.figure(figsize=(20,20))
sns.heatmap(df[top_features].corr(), annot=True)
```

```
[18]: <AxesSubplot:
```



Information Gain

```
[22]: from sklearn.feature_selection import mutual_info_classif
```

```
[23]: mutual_info=mutual_info_classif(X,y)
```

```
[24]: mutual_data=pd.Series(mutual_info, index=X.columns)
mutual_data.sort_values(ascending=False)
```

```
[24]: PriorDefault_t      0.297249
CreditScore            0.167713
Employed_t              0.108918
Income                  0.083886
EducationLevel          0.077240
YearsEmployed           0.073430
Ethnicity                0.051324
Married_u                0.029247
BankCustomer_gg          0.022781
Citizen_p                 0.020874
BankCustomer_P            0.020765
Debt                      0.017180
Married_y                 0.012308
Age                      0.007701
Gender_b                  0.000000
Citizen_s                  0.000000
Married_l                  0.000000
BankCustomer_g              0.000000
DriversLicense_t            0.000000
dtype: float64
```

```
[24]: Gender_b      0.000000
Citizen_s      0.000000
Married_l      0.000000
BankCustomer_g  0.000000
DriversLicense_t 0.000000
dtype: float64
```

```
[25]: ## Not including Employed because Income is same as Employed
df = df[['PriorDefault_t','YearsEmployed','CreditScore','Income','Approved']]
df.head()
```

```
[25]: PriorDefault_t  YearsEmployed  CreditScore  Income  Approved
0               1        1.25       1  321.692754      1
1               1         3.04       6  560.000000      1
2               1         1.50       0  824.000000      1
3               1         3.75       5  300.000000      1
4               1         1.71       0  321.692754      1
```

```
[26]: df.to_csv('feature_selection.csv',index=False)
```

Training with KNN

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

[2]: import warnings
warnings.filterwarnings('ignore')

[3]: df=pd.read_csv('feature_Selection.csv')
df.head()

[3]:   PriorDefault_t  YearsEmployed  CreditScore  Income  Approved
0            1           1.25          1  321.692754      1
1            1           3.04          6  560.000000      1
2            1           1.50          0  824.000000      1
3            1           3.75          5  3.000000      1
4            1           1.71          0  321.692754      1

[4]: from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
scaled_features = scaler.fit_transform(df.drop('Approved',axis=1))

[5]: df_feat = pd.DataFrame(scaled_features,columns=df.columns[:-1])
df_feat.head()

[5]:   PriorDefault_t  YearsEmployed  CreditScore  Income

```



```
[6]: ##### Train Test Split
from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test=train_test_split(scaled_features,df['Approved'],test_size=0.20,random_state=0)

[7]: y_test.shape
[7]: (138,)

[8]: y_test.value_counts()
[8]: 0    78
     1    60
Name: Approved, dtype: int64

[9]: y_train.value_counts()
[9]: 0    305
     1    247
Name: Approved, dtype: int64

[10]: from sklearn.neighbors import KNeighborsClassifier
knn_classifier=KNeighborsClassifier(n_neighbors=1).fit(X_train,y_train)
prediction=knn_classifier.predict(X_test)


[11]: from sklearn.metrics import confusion_matrix,classification_report,accuracy_score,roc_curve, roc_auc_score
print(f'confusion matrix : {confusion_matrix(y_test,prediction)}')
print(f'Accuracy Score : {accuracy_score(y_test,prediction)}')
print(classification_report(y_test,prediction))
confusion matrix : [[65 13]
 [15 45]]
Accuracy Score : 0.7971014402752623
```



```
Name: Approved, dtype: int64
[13]: auc = roc_auc_score(y_test, prediction)
auc
[13]: 0.7916666666666667

[14]: fpr, tpr, thresholds = roc_curve(y_test, prediction)

[15]: plt.plot(fpr, tpr, color='orange', label='ROC')
plt.plot([0, 1], [0, 1], color='darkblue', linestyle='--',label='ROC curve (area = %0.2f)' % auc)
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve for Naive Bayes')
plt.legend()
plt.show()
```



```

[16]: from sklearn.model_selection import RandomizedSearchCV
random_grid = { 'algorithm' : [ 'ball_tree', 'kd_tree', 'brute'],
                'leaf_size' : [18,20,25,27,30,32,34],
                'n_neighbors' : [3,5,7,9,10,11,12,13]
            }

[17]: # Use the random grid to search for best hyperparameters
# First create the base model to tune
Classifier=KNeighborsClassifier()

[18]: # Random search of parameters, using 5 fold cross validation,
# search across 100 different combinations
knn_random = RandomizedSearchCV(estimator = Classifier, param_distributions = random_grid, cv = 5, verbose=2)

[19]: knn_random.fit(X_train,y_train)
...
[20]: knn_random.best_params_
{'n_neighbors': 5, 'leaf_size': 34, 'algorithm': 'kd_tree'}

[21]: best_random_grid=knn_random.best_estimator_

[22]: from sklearn.metrics import accuracy_score,accuracy_score,roc_curve, roc_auc_score
y_pred=best_random_grid.predict(X_test)
print(confusion_matrix(y_test,y_pred))
print("Accuracy Score {}".format(accuracy_score(y_test,y_pred)))
print("Classification report: {}".format(classification_report(y_test,y_pred)))

[28]: knn = KNeighborsClassifier(n_neighbors=3)

knn.fit(X_train,y_train)
pred = knn.predict(X_test)

print(accuracy_score(y_test,pred))
print('\n')
print(confusion_matrix(y_test,pred))
print('\n')
print(classification_report(y_test,pred))

0.8913043478260869

[[85  1]
 [14 38]]


      precision    recall  f1-score   support

          0       0.86      0.99      0.92       86


[29]: # Area Under Curve
auc = roc_auc_score(y_test, pred)
auc

[29]: 0.8595706618962433

[30]: fpr, tpr, thresholds = roc_curve(y_test, pred)

[30]: fpr, tpr, thresholds = roc_curve(y_test, pred)

[31]: plt.plot(fpr, tpr, color='orange', label='ROC')
plt.plot([0, 1], [0, 1], color='darkblue', linestyle='--',label='ROC curve (area = %0.2f)' % auc)
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve for Naive Bayes')
plt.legend()
plt.show()


[32]: import pickle
# save the model to disk
filename = 'finalized_model_knn.sav'
pickle.dump(knn,open(filename,'wb'))

```

Training with Naive Bayes

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

[2]: import warnings
warnings.filterwarnings('ignore')

[3]: df=pd.read_csv('feature_Selection.csv')
df.head()

[3]:   PriorDefault_t  YearsEmployed  CreditScore  Income  Approved
0             1          1.25           1  321.692754      1
1             1          3.04           6  560.000000      1
2             1          1.50           0  824.000000      1
3             1          3.75           5  3.000000      1
4             1          1.71           0  321.692754      1

[4]: from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
scaled_features = scaler.fit_transform(df.drop('Approved',axis=1))

[5]: df_feat = pd.DataFrame(scaled_features,columns=df.columns[:-1])
df_feat.head()

[5]:   PriorDefault_t  YearsEmployed  CreditScore  Income

```



```
[6]: ##### Train Test Split
from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test=train_test_split(scaled_features,df['Approved'],test_size=0.20,random_state=10)

[7]: from sklearn.naive_bayes import GaussianNB
NB=GaussianNB().fit(X_train,y_train)
prediction=NB.predict(X_test)

[8]: from sklearn.metrics import confusion_matrix,classification_report,accuracy_score,roc_curve, roc_auc_score
print(f'confusion matrix : {confusion_matrix(y_test,prediction)}')
print(f'Accuracy Score : {accuracy_score(y_test,prediction)}')
print(classification_report(y_test,prediction))

confusion matrix : [[66  4]
 [20 48]]
Accuracy Score : 0.8260869565217391
      precision    recall  f1-score   support
          0       0.77     0.94     0.85      70
          1       0.92     0.71     0.80      68

      accuracy                           0.83      138
     macro avg       0.85     0.82     0.82      138
weighted avg       0.84     0.83     0.82      138

[9]: # Area Under Curve
auc = roc_auc_score(y_test, prediction)
auc

[9]: 0.8243697478991597
```



```
[10]: fpr, tpr, thresholds = roc_curve(y_test, prediction)

[11]: plt.plot(fpr, tpr, color='orange', label='ROC')
plt.plot([0, 1], [0, 1], color='darkblue', linestyle='--',label='ROC curve (area = %0.2f)' % auc)
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve for Naive Bayes')
plt.legend()
plt.show()
```

The figure shows a Receiver Operating Characteristic (ROC) curve plot. The x-axis is labeled 'False Positive Rate' and ranges from 0.0 to 1.0. The y-axis is labeled 'True Positive Rate' and ranges from 0.0 to 1.0. A solid orange line represents the ROC curve, starting at (0,0) and ending at (1,1). A dashed blue diagonal line represents the random classifier baseline. The area under the ROC curve is labeled as 0.82.


```
[12]: import pickle
# save the model to disk
filename = 'finalized_model_nb.sav'
pickle.dump(NB, open(filename, 'wb'))
```

Training with Decision Tree

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

[2]: import warnings
warnings.filterwarnings('ignore')

[3]: df=pd.read_csv('feature_Selection.csv')
df.head()

[3]:   PriorDefault YearsEmployed CreditScore Income Approved
0             1        1.25      1 321.692754      1
1             1        3.04      6 560.000000      1
2             1        1.50      0 824.000000      1
3             1        3.75      5  3.000000      1
4             1        1.71      0 321.692754      1

[4]: ##### Independent And Dependent features
X=df.drop('Approved',axis=1)
y=df['Approved']

[5]: ##### Train Test Split
from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.20,random_state=0)

[6]: from sklearn.tree import DecisionTreeClassifier

[6]: from sklearn.tree import DecisionTreeClassifier
df_classifier=DecisionTreeClassifier().fit(X_train,y_train)
prediction=df_classifier.predict(X_test)

[7]: from sklearn import tree

[8]: plt.figure(figsize=(30,10))
tree.plot_tree(df_classifier,filled=True)

[8]: [Text(670.425986421053, 525.48, 'X[0] <= 0.5\ngini = 0.494\nsamples = 552\nvalue = [305, 247]'),
Text(268.72105263157897, 489.24, 'X[1] <= 0.062\ngini = 0.123\nsamples = 259\nvalue = [242, 17]'),
Text(114.53684210526316, 453.0, 'X[3] <= 310.846\ngini = 0.245\nsamples = 56\nvalue = [48, 8]'),
Text(70.48421052631579, 416.76, 'X[3] <= 3.0\ngini = 0.074\nsamples = 26\nvalue = [25, 1]'),
Text(52.863157894736844, 380.52, 'X[2] <= 0.5\ngini = 0.245\nsamples = 7\nvalue = [6, 1]'),
Text(35.242105263157896, 344.28, 'X[1] <= 0.02\ngini = 0.444\nsamples = 3\nvalue = [2, 1]'),
Text(17.621052631578948, 308.04, 'gini = 0.5\nsamples = 2\nvalue = [1, 1]'),
Text(52.863157894736844, 308.04, 'gini = 0.0\nsamples = 1\nvalue = [1, 0]'),
Text(70.48421052631579, 344.28, 'gini = 0.0\nsamples = 4\nvalue = [4, 0]'),
Text(88.10526315789474, 380.52, 'gini = 0.0\nsamples = 19\nvalue = [19, 0]'),
Text(158.58947368421053, 416.76, 'X[3] <= 1347.0\ngini = 0.358\nsamples = 30\nvalue = [23, 7]'),
Text(123.34736842105264, 380.52, 'X[3] <= 336.346\ngini = 0.337\nsamples = 28\nvalue = [22, 6]'),
Text(105.72631578947369, 344.28, 'X[2] <= 0.5\ngini = 0.386\nsamples = 23\nvalue = [17, 6]'),
Text(88.10526315789474, 308.04, 'X[1] <= 0.02\ngini = 0.351\nsamples = 22\nvalue = [17, 5]'),
Text(70.48421052631579, 271.8, 'gini = 0.346\nsamples = 18\nvalue = [14, 4]'),
Text(105.72631578947369, 271.8, 'gini = 0.375\nsamples = 4\nvalue = [3, 1]'),
Text(123.34736842105264, 308.04, 'gini = 0.0\nsamples = 1\nvalue = [0, 1]'),
Text(140.96842105263158, 344.28, 'gini = 0.0\nsamples = 5\nvalue = [5, 0]'),
Text(193.83157894736843, 380.52, 'X[1] <= 0.02\ngini = 0.5\nsamples = 2\nvalue = [1, 1]'),
Text(176.21052631578948, 344.28, 'gini = 0.0\nsamples = 1\nvalue = [0, 1]'),
Text(211.45263157894738, 344.28, 'gini = 0.0\nsamples = 1\nvalue = [0, 1]'),
Text(143.20526315789473, 453.0, 'X[1] <= 5\ngini = 0.225\nsamples = 200\nvalue = [104, 96]')

[Text(1630.978547506721, 105.07555555555556, 'gini = 0.0\nsamples = 0\nvalue = [0, 0]'),
Text(1638.7578947368422, 380.52, 'X[3] <= 357.0\ngini = 0.02\nsamples = 101\nvalue = [1, 100]'),
Text(1621.1368421052632, 344.28, 'X[3] <= 337.0\ngini = 0.062\nsamples = 31\nvalue = [1, 30]'),
Text(1603.5157894736842, 308.04, 'gini = 0.0\nsamples = 30\nvalue = [0, 30]'),
Text(1638.7578947368422, 308.04, 'gini = 0.0\nsamples = 1\nvalue = [1, 0]'),
Text(1656.378947368421, 344.28, 'gini = 0.0\nsamples = 70\nvalue = [0, 70]')

<img alt="A complex decision tree diagram with many nodes and branches, colored in orange and blue. The tree has multiple levels of splits, with some nodes being terminal leaf nodes and others being internal nodes leading to further splits." data-bbox="195 685 955 860]

[9]: from sklearn.metrics import confusion_matrix,classification_report,accuracy_score
print(f'confusion matrix : {confusion_matrix(y_test,prediction)}')
print(f'Accuracy Score : {accuracy_score(y_test,prediction)}')</pre>

```

```
[9]: from sklearn.metrics import confusion_matrix,classification_report,accuracy_score
print(f'confusion matrix : {confusion_matrix(y_test,prediction)})')
print(f'Accuracy Score : {accuracy_score(y_test,prediction)})')
print(classification_report(y_test,prediction))
```

```
confusion matrix : [[67 11]
 [16 44]]
Accuracy Score : 0.8943478260869565
          precision    recall   f1-score   support
          0       0.81     0.86     0.83      78
          1       0.80     0.73     0.77      60

   accuracy           0.80      138
  macro avg       0.80     0.80     0.80      138
weighted avg       0.80     0.80     0.80      138
```

```
[10]: from sklearn.model_selection import RandomizedSearchCV
random_grid={
 "splitter" : ["best", "random"] ,
 "max_depth" : [ 3, 4, 5, 6, 8, 10, 12, 15],
 "min_samples_leaf" : [ 1,2,3,4,5 ],
 "min_weight_fraction_leaf": [0.1,0.2,0.3,0.4],
 "max_features" : ["auto","log2","sqrt",None ],
 "max_leaf_nodes": [None,10,20,30,40,50,60,70]
 }
```

```
}
```

```
[11]: # Use the random grid to search for best hyperparameters
# First create the base model to tune
Classifier=DecisionTreeClassifier()
```

```
[12]: # Random search of parameters, using 3 fold cross validation,
# search across 100 different combinations
df_random = RandomizedSearchCV(estimator = Classifier, param_distributions = random_grid, n_iter = 100, cv = 5, verbose=2, random_state=10, n_jobs = 1)
```

```
[13]: df_random.fit(X_train,y_train)
***
```

```
[14]: df_random.best_params_
```

```
{'splitter': 'random',
'min_weight_fraction_leaf': 0.3,
'min_samples_leaf': 4,
'max_leaf_nodes': None,
'max_features': None,
'max_depth': 6}
```

```
[15]: best_random_grid=df_random.best_estimator_
```

```
[16]: print(best_random_grid)
```

```
DecisionTreeClassifier(max_depth=6, min_samples_leaf=4,
min_weight_fraction_leaf=0.3, splitter='random')
```

```
[15]: best_random_grid=df_random.best_estimator_
```

```
[16]: print(best_random_grid)
```

```
DecisionTreeClassifier(max_depth=6, min_samples_leaf=4,
min_weight_fraction_leaf=0.3, splitter='random')
```

```
[17]: from sklearn.metrics import accuracy_score
y_pred=best_random_grid.predict(X_test)
print(confusion_matrix(y_test,y_pred))
print("Accuracy Score: {} ".format(accuracy_score(y_test,y_pred)))
print("Classification report: {} ".format(classification_report(y_test,y_pred)))
```

```
[[64 14]
```

```
 [ 6 54]]
Accuracy Score: 0.855072463768116
```

```
Classification report:          precision    recall   f1-score   support
```

	precision	recall	f1-score	support
0	0.91	0.82	0.86	78
1	0.79	0.90	0.84	60

```
accuracy           0.86      138
```

```
macro avg       0.85     0.86     0.85      138
```

```
weighted avg       0.86     0.86     0.86      138
```

```
[18]: import pickle
```

```
# save the model to disk
```

```
filename='finalized_model_DT.sav'
```

```
pickle.dump(df_random, open(filename, 'wb'))
```

Training with Random Forest and Hyper parameter tuning

All Techniques Of Hyper Parameter Optimization

```
1.GridSearchCV  
2.RandomizedSearchCV  
3.Bayesian Optimization -Automate Hyperparameter Tuning (Hyperopt)  
4.Sequential Model based optimization(tuning a scikit-learn estimator with skopt)  
5.Optuna- Automate Hyperparameter Tuning  
6.Genetic Algorithms (TPOT Classifier)
```

```
[1]: import warnings  
warnings.filterwarnings('ignore')  
  
[2]: import pandas as pd  
df=pd.read_csv('feature_Selection.csv')  
df.head()  
  
[2]: PriorDefault_t  YearsEmployed CreditScore Income Approved  
0 1 1.25 1 321.692754 1  
1 1 3.04 6 560.000000 1  
2 1 1.50 0 824.000000 1  
3 1 3.75 5 3.000000 1  
4 1 1.71 0 321.692754 1  
  
[3]: ##### Independent And Dependent features  
X=df.drop('Approved',axis=1)  
y=df['Approved']  
#all the code till here is same as previous notebook
```

```
[4]: ##### Train Test Split  
from sklearn.model_selection import train_test_split  
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.20,random_state=0)  
  
[5]: from sklearn.ensemble import RandomForestClassifier  
rf_classifier=RandomForestClassifier(n_estimators=10).fit(X_train,y_train)  
prediction=rf_classifier.predict(X_test)  
  
[6]: from sklearn.metrics import confusion_matrix,classification_report,accuracy_score  
print(f'confusion matrix : {confusion_matrix(y_test,prediction)}')  
print(f'Accuracy Score : {accuracy_score(y_test,prediction)}')  
print(classification_report(y_test,prediction))  
  
confusion matrix : [[67 11]  
 [14 46]]  
Accuracy Score : 0.8188405797101449  
precision recall f1-score support  
  
          0 0.83 0.86 0.84 78  
          1 0.81 0.77 0.79 60  
  
accuracy 0.82 0.82 0.82 138  
macro avg 0.82 0.81 0.81 138  
weighted avg 0.82 0.82 0.82 138
```

Manual Hyperparameter Tuning

```
[7]: ##### Manual Hyperparameter Tuning  
model=RandomForestClassifier(n_estimators=300,criterion='entropy',  
                             max_features='sqrt',min_samples_leaf=10,random_state=100).fit(X_train,y_train)  
predictions=model.predict(X_test)  
print(confusion_matrix(y_test,predictions))  
print(accuracy_score(y_test,predictions))  
print(classification_report(y_test,predictions))  
  
[[67 11]  
 [10 58]]  
0.8478260869565217  
precision recall f1-score support  
  
          0 0.87 0.86 0.86 78  
          1 0.82 0.83 0.83 60  
  
accuracy 0.85 0.85 0.85 138  
macro avg 0.84 0.85 0.85 138  
weighted avg 0.85 0.85 0.85 138
```

Randomized Search Cv

```
[8]: import numpy as np
from sklearn.model_selection import RandomizedSearchCV
# Number of trees in random forest
n_estimators = [int(x) for x in np.linspace(start = 200, stop = 2000, num = 10)]
# Number of features to consider at every split
max_features = ['auto', 'sqrt','log2']
# Maximum number of levels in tree
max_depth = [int(x) for x in np.linspace(10, 1000,10)]
# Minimum number of samples required to split a node
min_samples_split = [2, 5, 10,14]
# Minimum number of samples required at each leaf node
min_samples_leaf = [1, 2, 4,6,8]
# Create the random grid
random_grid = {'n_estimators': n_estimators,
               'max_features': max_features,
               'max_depth': max_depth,
               'min_samples_split': min_samples_split,
               'min_samples_leaf': min_samples_leaf,
               'criterion':['entropy','gini']}
print(random_grid)

{'n_estimators': [200, 400, 600, 800, 1000, 1200, 1400, 1600, 1800, 2000], 'max_features': ['auto', 'sqrt', 'log2'], 'max_depth': [10, 120, 230, 340, 450, 560, 670, 780, 890, 1000], 'min_samples_split': [2, 5, 10, 14], 'min_samples_leaf': [1, 2, 4, 6, 8], 'criterion': ['entropy', 'gini']}

[9]: rf=RandomForestClassifier()
rf_randomcv=RandomizedSearchCV(estimator=rf,param_distributions=random_grid,n_iter=100, cv=3,verbose=2,
                                random_state=100,n_jobs=-1)
```

```
[10]: rf_randomcv.best_params_
[10]: {'n_estimators': 1400,
       'min_samples_split': 10,
       'min_samples_leaf': 8,
       'max_features': 'auto',
       'max_depth': 450,
       'criterion': 'gini'}
[11]: rf_randomcv
[11]: RandomizedSearchCV(cv=3, estimator=RandomForestClassifier(), n_iter=100,
                        n_jobs=-1,
                        param_distributions={'criterion': ['entropy', 'gini'],
                                             'max_depth': [10, 120, 230, 340, 450,
                                                          560, 670, 780, 890,
                                                          1000],
                                             'max_features': ['auto', 'sqrt',
                                                               'log2'],
                                             'min_samples_leaf': [1, 2, 4, 6, 8],
                                             'min_samples_split': [2, 5, 10, 14],
                                             'n_estimators': [200, 400, 600, 800,
                                                             1000, 1200, 1400, 1600,
                                                             1800, 2000]},
                        random_state=100, verbose=2)
[12]: best_random_grid=rf_randomcv.best_estimator_
[13]: from sklearn.metrics import accuracy_score
y_pred=best_random_grid.predict(X_test)
print(confusion_matrix(y_test,y_pred))
[[[[Accuracy Score: 0.8623188405797102
Classification report: precision recall f1-score support
          0      0.87      0.88      0.88      78
          1      0.85      0.83      0.84      60
accuracy      0.86      0.86      0.86     138
macro avg      0.86      0.86      0.86     138
weighted avg      0.86      0.86      0.86     138]]]
```

GridSearch CV

```
[14]: from sklearn.model_selection import GridSearchCV
param_grid = {
    'criterion': [rf_randomcv.best_params_['criterion']],
    'max_depth': [rf_randomcv.best_params_['max_depth']],
    'max_features': [rf_randomcv.best_params_['max_features']],
    'min_samples_leaf': [rf_randomcv.best_params_['min_samples_leaf'],
                         rf_randomcv.best_params_['min_samples_leaf']+2,
                         rf_randomcv.best_params_['min_samples_leaf']+4],
    'min_samples_split': [rf_randomcv.best_params_['min_samples_split']+2,
                          rf_randomcv.best_params_['min_samples_split']+1,
                          rf_randomcv.best_params_['min_samples_split'],
                          rf_randomcv.best_params_['min_samples_split']+1,
                          rf_randomcv.best_params_['min_samples_split']+2],
```

```

[15]: ##### Fit the grid_search to the data
rf=RandomForestClassifier()
grid_search=GridSearchCV(estimator=rf,param_grid=param_grid, cv=10,n_jobs=-1,verbose=2)
grid_search.fit(X_train,y_train)

Fitting 10 folds for each of 75 candidates, totalling 750 fits
[15]: GridSearchCV(cv=10, estimator=RandomForestClassifier(), n_jobs=-1,
                  param_grid=[{'criterion': ['gini'], 'max_depth': [450],
                               'max_features': ['auto'],
                               'min_samples_leaf': [8, 10, 12],
                               'min_samples_split': [8, 9, 10, 11, 12],
                               'n_estimators': [1200, 1300, 1400, 1500, 1600]},
                               verbose=2)

[16]: grid_search.best_estimator_

[16]: RandomForestClassifier(max_depth=450, min_samples_leaf=10, min_samples_split=9,
                           n_estimators=1200)

[17]: best_grid=grid_search.best_estimator_

[18]: best_grid

[18]: RandomForestClassifier(max_depth=450, min_samples_leaf=10, min_samples_split=9,
                           n_estimators=1200)

[19]: y_pred=best_grid.predict(X_test)
print(confusion_matrix(y_test,y_pred))
print("Accuracy Score {}".format(accuracy_score(y_test,y_pred)))
print("Classification report: {}".format(classification_report(y_test,y_pred)))

[[68 10]]

```

Automated Hyperparameter Tuning

Automated Hyperparameter Tuning can be done by using techniques such as

- Bayesian Optimization
- Gradient Descent
- Evolutionary Algorithms

Bayesian Optimization

Bayesian optimization uses probability to find the minimum of a function. The final aim is to find the input value to a function which can give us the lowest possible output value. It usually performs better than random.grid and manual search providing better performance in the testing phase and reduced optimization time. In Hyperopt, Bayesian Optimization can be implemented giving 3 three main parameters to the function fmin.

- Objective Function = defines the loss function to minimize.
- Domain Space = defines the range of input values to test (in Bayesian Optimization this space creates a probability distribution for each of the used Hyperparameters).
- Optimization Algorithm = defines the search algorithm to use to select the best input values to use in each new iteration.

```

[20]: from hyperopt import hp,fmin,tpe,STATUS_OK,Trials

[21]: space = {'criterion': hp.choice('criterion', ['entropy', 'gini']),
            'max_depth': hp.choice('max_depth',[10, 120, 230, 340, 450, 560, 670, 780, 890, 1000]),
            'max_features': hp.choice('max_features', ['auto', 'sqrt','log2', None]),
            'min_samples_leaf': hp.uniform('min_samples_leaf', 0, 0.5),
            'min_samples_split' : hp.uniform ('min_samples_split', 0, 1),
            'n_estimators' : hp.choice('n_estimators', [200, 400, 600, 800, 1000, 1200, 1400, 1600, 1800, 2000])
            }

[22]: def objective(space):
        model = RandomForestClassifier(criterion = space['criterion'], max_depth = space['max_depth'],
                                       max_features = space['max_features'])

[23]: from sklearn.model_selection import cross_val_score
trials = Trials()
best = fmin(fn= objective,
            space= space,
            algo= tpe.suggest,
            max_evals = 80,
            trials= trials)
best

100%|██████████| 80/80 [08:56<00:00, 6.71s/trial, best loss: -0.8585913185913185]
[23]: {'criterion': 1,
       'max_depth': 0,
       'max_features': 1,
       'min_samples_leaf': 0.08453684676276399,
       'min_samples_split': 0.11006204582812613,
       'n_estimators': 2}

[24]: crit = {0: 'entropy', 1: 'gini'}
depth = [10, 120, 230, 340, 450, 560, 670, 780, 890, 1000]
feat = {0: 'auto', 1: 'sqrt', 2: 'log2', 3: None}
est = {0: 200, 1: 400, 2: 600, 3: 800, 4: 1000, 5:1200, 6:1400,7:1600,8:1800,9:2000}

print(crit[best['criterion']])
print(feat[best['max_features']])
print(est[best['n_estimators']])

gini
sqrt
600

[25]: best['min_samples_leaf']

```

```
[26]: trainedforest = RandomForestClassifier(criterion = crit[best['criterion']], max_depth = depth[best['max_depth']],
                                           max_features = feat[best['max_features']],
                                           min_samples_leaf = best['min_samples_leaf'],
                                           min_samples_split = best['min_samples_split'],
                                           n_estimators = est[best['n_estimators']]).fit(X_train,y_train)
predictionforest = trainedforest.predict(X_test)
print(confusion_matrix(y_test,predictionforest))
print(accuracy_score(y_test,predictionforest))
print(classification_report(y_test,predictionforest))
acc5 = accuracy_score(y_test,predictionforest)
```

[[64 14]	
[6 54]]	
0.855072463768116	
	precision recall f1-score support
0	0.91 0.82 0.86 78
1	0.79 0.90 0.84 60
	accuracy 0.86 138
	macro avg 0.85 0.86 0.85 138
	weighted avg 0.86 0.86 0.86 138

```
[27]: import numpy as np
from sklearn.model_selection import RandomizedSearchCV
# Number of trees in random forest
n_estimators = [int(x) for x in np.linspace(start = 200, stop = 2000, num = 10)]
# Number of features to consider at every split
max_features = ['auto', 'sqrt','log2']
# Maximum number of levels in tree
max_depth = [int(x) for x in np.linspace(10, 1000,10)]
# Minimum number of samples required to split a node
min_samples_split = [2, 5, 10,14]
# Minimum number of samples required at each Leaf node
min_samples_leaf = [1, 2, 4,6,8]
# Create the random grid
param = {'n_estimators': n_estimators,
         'max_features': max_features,
         'max_depth': max_depth,
         'min_samples_split': min_samples_split,
         'min_samples_leaf': min_samples_leaf,
         'criterion':['entropy','gini']}
print(param)

{'n_estimators': [200, 400, 600, 800, 1000, 1200, 1400, 1600, 1800, 2000], 'max_features': ['auto', 'sqrt', 'log2'], 'max_depth': [10, 120, 230, 340, 450, 560, 670, 780, 890, 1000], 'min_samples_split': [2, 5, 10, 14], 'min_samples_leaf': [1, 2, 4, 6, 8], 'criterion': ['entropy', 'gini']}
```

```
[28]: from tpot import TPOTClassifier

tpot_classifier = TPOTClassifier(generations= 5, population_size= 24, offspring_size= 12,
                                  verbosity= 2, early_stop= 12,
                                  config_dict={'sklearn.ensemble.RandomForestClassifier': param},
                                  cv = 4, scoring = 'accuracy')
```

```
TPOTClassifier(generations=5, max_eval_time=None, max_memory=None,
               max_time=None, min_eval_time=None, min_time=None,
               n_jobs=-1, population_size=24, scoring='accuracy', verbosity=2)
```

```
[29]: accuracy = tpot_classifier.score(X_test, y_test)
print(accuracy)

0.8478260869565217
```

Optimize hyperparameters of the model using Optuna

The hyperparameters of the above algorithm are `n_estimators` and `max_depth` for which we can try different values to see if the model accuracy can be improved. The objective function is modified to accept a trial object. This trial has several methods for sampling hyperparameters. We create a study to run the hyperparameter optimization and finally read the best hyperparameters.

```
[30]: import optuna
import sklearn.svm
def objective(trial):
    classifier = trial.suggest_categorical('classifier', ['RandomForest', 'SVC'])

    if classifier == 'RandomForest':
        n_estimators = trial.suggest_int('n_estimators', 200, 2000,10)
        max_depth = int(trial.suggest_float('max_depth', 10, 100, log=True))

        clf = sklearn.ensemble.RandomForestClassifier(
            n_estimators=n_estimators, max_depth=max_depth)
    else:
        c = trial.suggest_float('svc_c', 1e-10, 1e10, log=True)
        clf = sklearn.svm.SVC(C=c, gamma='auto')
    return sklearn.model_selection.cross_val_score(
        clf,X_train,y_train, n_jobs=-1, cv=3).mean()
```

```

[31]: study = optuna.create_study(direction='maximize')
study.optimize(objective, n_trials=100)

trial = study.best_trial

print('Accuracy: {}'.format(trial.value))
print("Best hyperparameters: {}".format(trial.params))

[I 2021-11-07 09:21:53,119] A new study created in memory with name: no-name-d1f2e607-527d-4c27-ad91-968e601eb6d4
[I 2021-11-07 09:21:58,012] Trial 0 finished with value: 0.8278985507246377 and parameters: {'classifier': 'RandomForest', 'n_estimators': 1860, 'max_depth': 52.831636492821204}. Best is trial 0 with value: 0.8278985507246377.
[I 2021-11-07 09:22:03,753] Trial 1 finished with value: 0.6721014492753623 and parameters: {'classifier': 'SVC', 'svc_c': 274357.6778465932}. Best is trial 0 with value: 0.8278985507246377.
[I 2021-11-07 09:22:06,332] Trial 2 finished with value: 0.8333333333333334 and parameters: {'classifier': 'RandomForest', 'n_estimators': 700, 'max_depth': 66.28649973761492}. Best is trial 2 with value: 0.8333333333333334.
[I 2021-11-07 09:22:07,029] Trial 3 finished with value: 0.8278985507246377 and parameters: {'classifier': 'RandomForest', 'n_estimators': 420, 'max_depth': 21.09713286033926}. Best is trial 2 with value: 0.8333333333333334.
[I 2021-11-07 09:25:45,166] Trial 4 finished with value: 0.5887681159420289 and parameters: {'classifier': 'SVC', 'svc_c': 5861527765.6813545}. Best is trial 0 with value: 0.8278985507246377.
[I 2021-11-07 09:25:45,223] Trial 5 finished with value: 0.7101449275362318 and parameters: {'classifier': 'SVC', 'svc_c': 37.30576591726301}. Best is trial 2 with value: 0.8333333333333334.
[I 2021-11-07 09:25:47,261] Trial 6 finished with value: 0.552536231884058 and parameters: {'classifier': 'SVC', 'svc_c': 2.091714942429609e-08}. Best is trial 2 with value: 0.8333333333333334.
[I 2021-11-07 09:25:47,807] Trial 7 finished with value: 0.8315217391384347 and parameters: {'classifier': 'RandomForest', 'n_estimators': 1620, 'max_depth': 17.311885144567185}. Best is trial 2 with value: 0.8333333333333334.
[I 2021-11-07 09:25:50,104] Trial 8 finished with value: 0.8278985507246377 and parameters: {'classifier': 'RandomForest', 'n_estimators': 1450, 'max_depth': 10.25498292302449}. Best is trial 2 with value: 0.8333333333333334.
[I 2021-11-07 09:25:50,738] Trial 9 finished with value: 0.8333333333333334 and parameters: {'classifier': 'RandomForest', 'n_estimators': 380, 'max_depth': 98.70403999854032}. Best is trial 2 with value: 0.8333333333333334.
[I 2021-11-07 09:25:52,429] Trial 10 finished with value: 0.8278985507246377 and parameters: {'classifier': 'RandomForest', 'n_estimators': 1000, 'max_depth': 54.859407243484796}. Best is trial 2 with value: 0.8333333333333334.

Accuracy: 0.8333333333333334
Best hyperparameters: {'classifier': 'RandomForest', 'n_estimators': 700, 'max_depth': 66.28649973761492}

[32]: trial

[32]: FrozenTrial(number=2, values=[0.8333333333333334], datetime_start=datetime.datetime(2021, 11, 7, 9, 22, 3, 754986), datetime_complete=datetime.datetime(2021, 11, 7, 9, 22, 6, 331670), params={'classifier': 'RandomForest', 'n_estimators': 700, 'max_depth': 66.28649973761492}, distributions={'classifier': CategoricalDistribution(choices=('RandomForest', 'SVC')), 'n_estimators': IntUniformDistribution(high=2000, low=200, step=10), 'max_depth': LogUniformDistribution(high=100.0, low=10.0)}, user_attrs={}, system_attrs={}, intermediate_values={}, trial_id=2, state=TrialState.COMPLETED, value=None)

[33]: study.best_params

[33]: {'classifier': 'RandomForest',
       'n_estimators': 700,
       'max_depth': 66.28649973761492}

[38]: rfc=RandomForestClassifier(n_estimators=study.best_params['n_estimators'],max_depth=study.best_params['max_depth'])
rfc.fit(X_train,y_train)
[[66 12]
 [14 46]]
0.8115942028985508
      precision    recall   f1-score   support
          0        0.82      0.85      0.84       78
          1        0.79      0.77      0.78       60

      accuracy                           0.81      138
     macro avg        0.81      0.81      0.81      138
  weighted avg        0.81      0.81      0.81      138

[39]: y_pred=rfc.predict(X_test)
print(confusion_matrix(y_test,y_pred))
print(accuracy_score(y_test,y_pred))
print(classification_report(y_test,y_pred))
[[66 12]
 [14 46]]
0.8115942028985508
      precision    recall   f1-score   support
          0        0.82      0.85      0.84       78
          1        0.79      0.77      0.78       60

      accuracy                           0.81      138
     macro avg        0.81      0.81      0.81      138
  weighted avg        0.81      0.81      0.81      138

[40]: import pickle
# save the model to disk
filename = 'finalized_model_RF.sav'
pickle.dump(rfc_randomcv, open(filename, 'wb'))

```

Training with Xgboost

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

[2]: import warnings
warnings.filterwarnings('ignore')

[3]: import pandas as pd
df=pd.read_csv('feature_Selection.csv')
df.head()

[3]:   PriorDefault_t  YearsEmployed  CreditScore  Income  Approved
0             1          1.25        1  321.692754      1
1             1          3.04        6  560.000000      1
2             1          1.50        0  824.000000      1
3             1          3.75        5  3.000000      1
4             1          1.71        0  321.692754      1

[4]: ##### Independent And Dependent features
X=df.drop('Approved',axis=1)
y=df['Approved']

[5]: ##### Train Test Split
from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.20,random_state=0)
```

```
[9]: import xgboost as xgb
xgb_classifier=xgb.XGBClassifier().fit(X_train,y_train)
prediction=xgb_classifier.predict(X_test)

[16:25:50] WARNING: C:/Users/Administrator/workspace/xgboost-win64_release_1.4.0/src/learner.cc:1095: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

[10]: from sklearn.metrics import confusion_matrix,classification_report,accuracy_score
print(f'confusion matrix : {confusion_matrix(y_test,prediction)}')
print(f'Accuracy Score : {accuracy_score(y_test,prediction)}')
print(classification_report(y_test,prediction))

confusion matrix : [[66 12]
 [14 46]]
Accuracy Score : 0.8115942028985508
              precision    recall   f1-score   support
0            0.82     0.85     0.84      78
1            0.79     0.77     0.78      60

       accuracy           0.81      138
      macro avg     0.81     0.81      138
weighted avg     0.81     0.81     0.81      138
```

```
[11]: ##### Manual Hyperparameter Tuning
model=xgb.XGBClassifier(n_estimators=300,learning_rate =0.05,
                         max_depth=5,subsample=0.7,min_child_weight=3).fit(X_train,y_train)
predictions=model.predict(X_test)
print(confusion_matrix(y_test,predictions))
print(accuracy_score(y_test,predictions))
print(classification_report(y_test,predictions))

[16:25:50] WARNING: C:/Users/Administrator/workspace/xgboost-win64_release_1.4.0/src/learner.cc:1095: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.
[[69  9]
 [13 47]]
0.8405797101449275
              precision    recall   f1-score   support
0            0.84     0.88     0.86      78
1            0.84     0.78     0.81      60

       accuracy           0.84      138
      macro avg     0.84     0.83     0.84      138
weighted avg     0.84     0.84     0.84      138
```

Randomized Search CV

```
[12]: from sklearn.model_selection import RandomizedSearchCV
# Number of trees in random forest
n_estimators = [int(x) for x in np.linspace(start = 100, stop = 1200, num = 12)]
# Various Learning rate parameters
```

```

[13]: # Use the random grid to search for best hyperparameters
# First create the base model to tune
Classifier=xgb.XGBClassifier()

[14]: # Random search of parameters, using 5 fold cross validation,
# search across 100 different combinations
xg_random = RandomizedSearchCV(estimator = Classifier, param_distributions = random_grid,scoring='roc_auc', n_iter = 100, cv = 5, verbose=2, random_state=42)

[15]: xg_random.fit(X_train,y_train)
"""

[17]: xg_random.best_params_
{'subsample': 0.6,
 'n_estimators': 100,
 'min_child_weight': 6,
 'max_depth': 15,
 'learning_rate': '0.05'}

[18]: best_random_grid=xg_random.best_estimator_

[19]: from sklearn.metrics import accuracy_score
y_pred=best_random_grid.predict(X_test)
print(confusion_matrix(y_test,y_pred))
print("Accuracy Score {}".format(accuracy_score(y_test,y_pred)))
print("Classification report: {}".format(classification_report(y_test,y_pred)))
[[69  9]
 [11 49]]
Accuracy Score 0.855072463768116

```

Bayesian Optimization

```

[20]: from hyperopt import hp,fmin,tpe,STATUS_OK,Trials

[21]: space = {
    'max_depth' : hp.choice('max_depth', range(5, 30, 1)),
    'learning_rate' : hp.quniform('learning_rate', 0.01, 0.5, 0.01),
    'n_estimators' : hp.choice('n_estimators', range(20, 205, 5)),
    'gamma' : hp.quniform('gamma', 0, 0.50, 0.01),
    'min_child_weight' : hp.quniform('min_child_weight', 1, 10, 1),
    'subsample' : hp.quniform('subsample', 0.1, 1, 0.01),
    'colsample_bytree' : hp.quniform('colsample_bytree', 0.1, 1.0, 0.01) }

[22]: def objective(space):
    import xgboost as xgb
    model = xgb.XGBClassifier(n_estimators = space['n_estimators'],
                               max_depth = int(space['max_depth']),
                               learning_rate = space['learning_rate'],
                               gamma = space['gamma'],
                               min_child_weight = space['min_child_weight'],
                               subsample = space['subsample'],
                               colsample_bytree = space['colsample_bytree'])

    accuracy = cross_val_score(model, X_train, y_train, cv = 5).mean()


```

```

[27]: from tpot import TPOTClassifier
from xgboost import XGBClassifier

tpot_classifier = TPOTClassifier(generations= 5, population_size= 24, offspring_size= 12,
                                 verbosity= 2, early_stop= 12,
                                 cv = 4, scoring = 'accuracy')
tpot_classifier.fit(X_train,y_train)

Optimization Progress:  0%|          | 0/84 [00:00<?, ?pipeline/s]

Generation 1 - Current best internal CV score: 0.8605072463768115
Generation 2 - Current best internal CV score: 0.8605072463768115
Generation 3 - Current best internal CV score: 0.8641304347826086
Generation 4 - Current best internal CV score: 0.8641304347826086
Generation 5 - Current best internal CV score: 0.8641304347826086

Best pipeline: ExtraTreesClassifier(VarianceThreshold(input_matrix, threshold=0.001), bootstrap=False, criterion=entropy, max_features=0.1, min_samples_leaf=5, min_samples_split=4, n_estimators=100)
[27]: TPOTClassifier(cv=4, early_stop=12, generations=5, offspring_size=12,
                  population_size=24, scoring='accuracy', verbosity=2)

[28]: accuracy = tpot_classifier.score(X_test, y_test)
print(accuracy)
0.855072463768116

[29]: import pickle

```

Comparing Bag of all 5 models vs KNN

```
[1]: import pickle
models = []
models.append(pickle.load(open('finalized_model_nb.sav','rb')))
models.append(pickle.load(open('finalized_model_knn.sav','rb')))
models.append(pickle.load(open('finalized_model_DT.sav','rb')))
models.append(pickle.load(open('finalized_model_RF.sav','rb')))
models.append(pickle.load(open('finalized_model_xgb.sav','rb')))

[2]: import numpy as np
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()

[3]: feat = np.array([[0,1.25,0,322]])

[4]: preds = []
for model in models:
    preds.append(model.predict(feat))

[5]: preds
[5]: [array([1], dtype=int64),
array([1], dtype=int64),
array([0], dtype=int64),
array([0], dtype=int64),
array([0], dtype=int64)]

[6]: for model in models:
    print(model)
GaussianNB()

[7]: def predict(feat,scaled_feat):
    preds = []
    i=0
    for model in models:
        if i<2:
            preds.append(model.predict(scaled_feat))
        else:
            preds.append(model.predict(feat))
        i+=1
    o = preds.count(1)
    z = preds.count(0)
    if(o>z):
        return 1
    else:
        return 0

[8]: import pandas as pd
df = pd.read_csv("feature_selection.csv")

[9]: X = df.drop(['Approved'],axis=1)
scaled_features = scaler.fit_transform(df.drop('Approved',axis=1))
y = df['Approved']

[10]: df_feat = pd.DataFrame(scaled_features,columns=df.columns[:-1])
df_feat.head()
[10]: PriorDefault YearsEmployed CreditScore Income
0      0.95465     -0.300502   -0.324711 -0.288910

[11]: pred = []
print(X.shape[0])
for i in range(X.shape[0]):
    feat = np.array(X.iloc[i].values)
    scaled_feat = np.array([df_feat.iloc[i].values])
    pred.append(predict(feat,scaled_feat))
690

[12]: from sklearn.metrics import confusion_matrix,classification_report,accuracy_score,roc_curve, roc_auc_score

[13]: accuracy_score(y,np.array(pred))
[13]: 0.8724637681159421

[14]: confusion_matrix(y,np.array(pred))

[14]: array([[331,  52],
       [ 36, 271]], dtype=int64)

[15]: print(classification_report(y,np.array(pred)))
          precision    recall  f1-score   support

              0       0.90      0.86      0.88      383
              1       0.84      0.88      0.86      307

  accuracy                           0.87      690
 macro avg       0.87      0.87      0.87      690
weighted avg       0.87      0.87      0.87      690
```

```

[16]: model = pickle.load(open('finalized_model_knn.sav','rb'))

[17]: pred=[]
    for i in range(X.shape[0]):
        scaled_feat = np.array([df_feat.iloc[i].values])
        pred.append(model.predict(scaled_feat))

[18]: accuracy_score(y,np.array(pred))

[18]: 0.8826086956521739

[19]: confusion_matrix(y,np.array(pred))

[19]: array([[351,  32],
           [ 49, 258]], dtype=int64)

[20]: print(classification_report(y,np.array(pred)))
      precision    recall  f1-score   support

          0       0.88      0.92      0.90      383
          1       0.89      0.84      0.86      307

   accuracy                           0.88      690
  macro avg       0.88      0.88      0.88      690
weighted avg       0.88      0.88      0.88      690

```

Therefore KNN with 3 neighbours has given best accuracy than bag of models.

Saving mean and standard deviations to use them for scaling

```

[1]: import pandas as pd

[2]: df = pd.read_csv("feature_selection.csv")

[3]: df.head()

[3]:
  PriorDefault_t  YearsEmployed  CreditScore  Income  Approved
0             1            1.25         1  321.692754      1
1             1            3.04         6  560.000000      1
2             1            1.50         0  824.000000      1
3             1            3.75         5  3.000000      1
4             1            1.71         0  321.692754      1

[4]: meanstd = pd.DataFrame([[df['PriorDefault_t'].mean(),df['PriorDefault_t'].std(),
                            [df['YearsEmployed'].mean(),df['YearsEmployed'].std()],
                            [df['CreditScore'].mean(),df['CreditScore'].std()],
                            [df['Income'].mean(),df['Income'].std()]],columns=["mean", "std"],
                           index=["PriorDefault_t","YearsEmployed","CreditScore","Income"]])

[5]: meanstd

[5]:
      mean      std
PriorDefault_t  0.523188  0.499824
YearsEmployed   2.040920  2.633902
CreditScore     2.156522  3.564275
Income          459.228061 476.394710

[6]: meanstd.reset_index(inplace=True)

[7]: meanstd.to_csv("meanstd.csv",index=False)

[8]: data = pd.read_csv("meanstd.csv")

[9]: data.head()

[9]:
      index      mean      std
0  PriorDefault_t  0.523188  0.499824
1  YearsEmployed   2.040920  2.633902
2  CreditScore     2.156522  3.564275
3    Income      459.228061 476.394710

[10]: data = pd.read_csv("meanstd.csv",index_col="index")

```

7.2 Testing

app.py

```
import numpy as np
import pandas as pd
from flask import Flask, request, jsonify, render_template
from flask_cors import CORS,cross_origin
import pickle

app = Flask(__name__)
meanstd = pd.read_csv("meanstd.csv",index_col="index")
model = pickle.load(open('finalized_model_knn.sav','rb'))

@app.route('/', methods=['GET'])
@cross_origin()
def index():
    return render_template('index.html')

def scale(lis):
    scaled = []
    ans = (lis[0]-meanstd.loc['PriorDefault_t']['mean'])/meanstd.loc['PriorDefault_t']['std']
    scaled.append(ans)
    ans = (lis[1]-meanstd.loc['YearsEmployed']['mean'])/meanstd.loc['YearsEmployed']['std']
    scaled.append(ans)
    ans = (lis[2]-meanstd.loc['CreditScore']['mean'])/meanstd.loc['CreditScore']['std']
    scaled.append(ans)
    ans = (lis[3]-meanstd.loc['Income']['mean'])/meanstd.loc['Income']['std']
    scaled.append(ans)
    # print(scaled)
    return scaled

@app.route('/predict', methods=['POST'])
@cross_origin()
def predict():
    if request.method == 'POST':
        try:
            features = [float(x) for x in request.form.values()]
            scaled_feat = scale(features)
            final_feat = [np.array(scaled_feat)]
            prediction = model.predict(final_feat)
            # print(prediction)
            if prediction==1:
                return render_template('index.html', prediction_text = "Credit Card Approved 😊")
            else:
                return render_template('index.html', prediction_text = "Credit Card Not Approved 😞")
        except Exception as e:
            print('The Exception message is: ',e)
            return 'something is wrong'
    else:
        return render_template('index.html')

if __name__ == "__main__":
    app.run(debug=True)
```

HTML and CSS

```
<!DOCTYPE html>
<html>
<head>
    <meta charset="UTF-8">
    <title> Credit Card Approval Prediction </title>
    <link rel="stylesheet" href="static/css/cssfile.css">
</head>
<body>
    <div class="login">
        <center><h1> Credit Card Approval Prediction </h1></center>
        <center>
            <h2>
                Welcome to my website.
                <br>
                <br>
                Enter the following data to predict the approval status of credit card.
            </h2>
        </center>
        <br><br>
        <form action="{{url_for('predict')}}" method="post" class="form-container">
            <div class="div">
                <label>
                    Prior Default 🤔 &nbsp;&nbsp; &nbsp;&nbsp;&nbsp;&nbsp;&nbsp;
                </label>
                <label>
                    <input type="number" min=0 max=2000 step=0.01 name="income" required="required">
                </label>
            </div>
            <br>
            <br>
            <div class="div">
                <button type="submit" class="btn btn-primary btn-block btn-large">
                    PREDICT
                </button>
            </div>
            <br>
            <br>
            <div class="q">
                <strong style="font-size:35px">{{ prediction_text }}</strong>
            </div>
        </form>
        <br>
        <br>
    </div>
</body>
</html>
```

```
body {}  
background-image: url('h2.jpg');  
height: 100%;  
background-position: center;  
background-repeat: no-repeat;  
background-size: cover;  
}  
input {  
width: 300px;  
height: 30px;  
font-size: 20px;  
}  
select {  
width: 308px;  
height: 40px;  
font-size: 20px;  
}  
label {  
font-size: 23px;  
color: #000000;  
}  
button {  
width: 200px;  
height: 30px;
```

```
button {  
width: 200px;  
height: 30px;  
font-size: 15px;  
}  
.form-container {  
width: 50%;  
}  
.div {  
display: flex;  
align-items: center;  
justify-content: space-around;  
}  
.q {  
display: flex;  
align-items: center;  
justify-content: space-around;  
}
```

Testing in Local Host

Credit Card Approval Prediction

Welcome to my website.


Enter the following data to predict the approval status of credit card.

Prior Default 😞

Years Employed 🧑‍🤝‍🧑

Credit Score 💹

Income 💰



Prior Default 😞

Years Employed 🧑‍🤝‍🧑

Credit Score 💹

Income 💰

Credit Card Approved 🎉

Prior Default 😞

Years Employed 🧑‍🤝‍🧑

Credit Score 📈

Income 💰

NO


2

2

700

PREDICT

Credit Card Not Approved 😞



Deployed web page

Credit Card Approval Prediction

Welcome to my website.

Enter the following data to predict the approval status of credit card.

Prior Default 😞

Years Employed 🧑‍🤝‍🧑

Credit Score 📈

Income 💰


YES

2

2

700

PREDICT



8. Results and Discussions

The accuracy of predicting the credit card is approved and credit card is not approved are almost the same 0.88 and 0.89. The classifier conservatively predicts that the status of credit card approval which will decrease a lot of work for the banks.

Performance measures of proposed algorithm

Accuracy Measure	Value
Specificity	0.91
Recall	0.84
Precision	0.88
F-Measure	0.86
Accuracy	0.88

The proposed algorithm returned an accuracy of 88.2% with performance measures: specificity, recall, precision and f-measure with values 0.91, 0.84, 0.88 and 0.86 respectively.

9. Conclusion and Future Scope

As of now the project has the models like KNN, Naive Bayes, Decision Tree, Random Forest, Xgboost. Neural networks can be used to get a better accuracy model. Neural network models can lead to give accuracies close to 98%.

10. References

- A) <https://www.ijrar.org/papers/IJRAR190B030.pdf>
- B) [https://www.researchgate.net/publication/321002603 Credit Approval Analysis using R](https://www.researchgate.net/publication/321002603_Credit_Approval_Analysis_using_R)
- C) <https://www.ijeat.org/wp-content/uploads/papers/v9i4/D7293049420.pdf>
- D) https://rstudio-pubs-static.s3.amazonaws.com/73039_9946de135c0a49daa7a0a9ed_a4a67a72.html
- E) <https://www.kaggle.com/echo9k/uci-credit-approval-data-set>
- F) <https://archive.ics.uci.edu/ml/datasets/Credit+Approval>